

Name of the student:	<i>Yash Sankpal</i>	Roll No.	<i>8695</i>
Practical Number:	4	Date of Practical:	
Relevant CO's	At the end of the course students will be able to use tools like hadoop and NoSQL to solve big data related problems.		
Sign here to indicate that you have read all the relevant material provided before attempting this practical			Sign:

Practical grading using Rubrics

Indicator	Very Poor	Poor	Average	Good	Excellent
Timeline (2)	More than a session late (0)	NA	NA	NA	Early or on time (2)
Code de- sign (2)	N/A	Very poor code design with no comments and indentation(0.5)	Poor code design with very comments and indentation (1)	Design with good coding standards (1.5)	Accurate design with better coding standards (2)
Performance (4)	Unable to perform the experiment (0)	Able to partially perform the experiment (1)	Able to perform the experiment for certain use cases (2)	Able to perform the experiment considering most of the use cases (3)	Able to perform the experiment considering all use cases (4)
Postlab (2)	No Execution(0)	N/A	Partially Executed (1)	N/A	Fully Executed (2)

Total Marks (10)	Sign of instructor with date

Practical

COURSE TITLE: BIG DATA ANALYTICS

COURSE TERM: 2021-2022

INSTRUCTOR NAME: SAURABH KULKARNI

Problem Statement: Perform matrix multiplication using one step map-reduce

Theory: Explain the concept of matrix multiplication using one step map-reduce with the help of an example

*For matrix multiplication, the number of columns in the first matrix must be equal to the number of rows in the second matrix. Thus, the matrix dimensions can be given as $A: i * j$ and $B: j * k$.*

For applying map-reduce paradigm to matrix multiplication, the rows from matrix A need to be grouped with columns of matrix B.

Thus, in mapper, we form key-value pairs as $(i, k) \rightarrow (A, j, val)$ where for matrix A, k loops in range of number of columns in B and for matrix B, k loops in range of number of rows in A.

Thus, for matrix A, we get pair i, k where for each row i , the reducer can work with k columns of B and for matrix B, we get pair i, k to correspond with row-wise pairs for A.

For eg:

A: 1 3 B: 4 5

4 6 7 8

After mapper, we get:

A: (1,1) -> (A, 1, 1)

(1,1) -> (A, 2, 3)

(2, 1) -> (A, 1, 4)

(2, 1) -> (A, 2, 6)

(1, 2) -> (A, 1, 1)

(1, 2) -> (A, 2, 3)

(2, 2) -> (A, 1, 4)

(2, 2) -> (A, 2, 6)

B: (1,1) -> (B, 1, 4)

(1, 1) -> (B, 2, 7)

(1,2) -> (B, 1, 5)

(1, 2) -> (B, 2, 8)

(2, 1) -> (B, 1, 4)

(2, 1) -> (B, 2, 7)

(2, 2) -> (B, 1, 5)

(2, 2) -> (B, 2, 8)

In reducer, we make:

(1, 1) -> Alist: [[A, 1, 1], [A, 2, 3]]

Blist: [[B, 1, 4], [B, 2, 7]]

...(for all keys)

Here, we multiply values with same j value and sum all products to get result matrix value for location (i, k)

Code:

code for mapper:

Code for Reducer:

Code for Driver Class:

```

1 package exp4;
2
3 import java.io.IOException;
4
5 import org.apache.hadoop.conf.Configuration;
6 import org.apache.hadoop.io.LongWritable;
7 import org.apache.hadoop.io.Text;
8 import org.apache.hadoop.mapreduce.Mapper;
9
10 public class Mapper4 extends Mapper<LongWritable, Text, Text, Text> {
11
12     public void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
13         Configuration conf = context.getConfiguration();
14
15         int m = Integer.parseInt(conf.get("m"));
16         int p = Integer.parseInt(conf.get("p"));
17
18         String str = value.toString();
19         String str[] = str.split(",");
20
21         Text outputKey = new Text();
22         Text outputValue = new Text();
23
24         if(str[0] == "A") {
25             for(int i=1; i<str.length(); i++) {
26                 //set outputkey as tokens[i], i.e. i,k
27                 outputKey.set(str[i], i);
28                 //set outputvalue as tokens[i], tokens[i] i.e. A,j,mj
29                 outputValue.set(str[i], str[i] + "," + str[i]);
30                 context.write(outputKey, outputValue);
31             }
32         }
33         else {
34             for(int i=0; i<str.length(); i++) {
35                 //set outputkey as i, tokens[i] i.e. i,k
36                 outputKey.set(i, str[i]);
37                 //set outputvalue as tokens[i], tokens[i] i.e. B,j,nk
38                 outputValue.set(str[i], str[i] + "," + str[i]);
39                 context.write(outputKey, outputValue);
40             }
41         }
42     }
43 }
44
45
46

```

```

1 package exp4;
2
3 import java.io.IOException;
4 import org.apache.hadoop.conf.Configuration;
5 import org.apache.hadoop.io.LongWritable;
6 import org.apache.hadoop.io.Text;
7 import org.apache.hadoop.mapreduce.Reducer;
8
9 public class Reducer4 extends Reducer<Text, Text, Text, Text> {
10
11     public void reduce(Text key, Iterable<Text> values, Context context) throws IOException, InterruptedException {
12         // process values
13         // process values
14         // define value as string type array
15         String[] value;
16
17         //create 2 hasmaps (hashA, hashB) of Integer,Float to store values from A and B matrix
18         HashMap<Integer, Float> aStore = new HashMap<Integer, Float>();
19         HashMap<Integer, Float> bStore = new HashMap<Integer, Float>();
20
21         //for each value in values
22         for(Each i: values) {
23             //convert each value to string as it is Text and split it by comma. Store this in value i.e. string type array defined above.
24             value = i.toString().split(",");
25
26             //if value[0] is A then
27             if(value[0].equals("A")) {
28                 //store value[1] as key and value[2] i.e. nk in hashA
29                 aStore.put(Integer.parseInt(value[1]), Float.parseFloat(value[2]));
30             }
31             else {
32                 //store value[1] as key and value[2] i.e. nk in hashB
33                 bStore.put(Integer.parseInt(value[1]), Float.parseFloat(value[2]));
34             }
35         }
36
37         Configuration conf = context.getConfiguration();
38
39         //take value of m which is common in both the matrices here from context object
40         int m = Integer.parseInt(conf.get("m"));
41
42         //define result of type float and assign value 0 to it.
43         float result = 0.0f;
44
45         //define a_i of type float
46         float a_i;
47
48         //define b_jk of type float
49         float b_jk;
50
51         //define a loop variable j and iterate till it is less than m
52         for(int j=0; j<m; j++) {
53             //check if value exists for a key j in hashA. If yes assign it to a_i else assign 0 to a_i.
54             a_i = aStore.containsKey(j) ? aStore.get(j).floatValue() : 0.0f;
55             //check if value exists for a key j in hashB. If yes assign it to b_jk else assign 0 to b_jk.
56             b_jk = bStore.containsKey(j) ? bStore.get(j).floatValue() : 0.0f;
57             //multiply a_i with b_jk and add this product to result which is of type float and declared above.
58             result += a_i * b_jk;
59
60         }
61
62         //if value of result is not 0 then
63         if(result != 0.0f) {
64             //write key,value to context, key is null and value is (i,k,result)
65             context.write(null, new Text(key.toString() + "," + Float.toString(result)));
66         }
67     }
68 }
69
70
71
72
73
74
75

```

```

1 package exp4;
2
3 import org.apache.hadoop.conf.Configuration;
4 import org.apache.hadoop.fs.Path;
5 import org.apache.hadoop.io.Text;
6 import org.apache.hadoop.mapreduce.Job;
7 import org.apache.hadoop.mapreduce.Mapper;
8 import org.apache.hadoop.mapreduce.Reducer;
9 import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
10 import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
11
12 public class Driver4 {
13
14     public static void main(String[] args) throws Exception {
15         Configuration conf = new Configuration();
16         conf.set("m", "3");
17         conf.set("n", "3");
18         conf.set("p", "3");
19         Job job = Job.getInstance(conf, "MatrixMultiplication");
20         job.setJarByClass(Driver4.class);
21         // TODO: specify a mapper
22         job.setMapperClass(Mapper4.class);
23         // TODO: specify a reducer
24         job.setReducerClass(Reducer4.class);
25
26         // TODO: specify output types
27         job.setOutputKeyClass(Text.class);
28         job.setOutputValueClass(Text.class);
29
30         // TODO: specify input and output DIRECTORIES (not files)
31         FileInputFormat.setInputPaths(job, new Path("hdfs://localhost:9000/input/exp4_postlab"));
32         FileOutputFormat.setOutputPath(job, new Path("hdfs://localhost:9000/output/exp4_"));
33
34         System.out.println(job.waitForCompletion(true));
35     }
36 }
37
38

```

PostLab:

1. Generate a 100x100 matrix in the format that above map-reduce code understands using suitable programming language
2. Compute execution time for a 100x100 matrix using suitable APIs

Code for postlab question

```
1 package exp4_postlab;
2
3 import org.apache.hadoop.conf.Configuration;
4 import org.apache.hadoop.fs.Path;
5 import org.apache.hadoop.io.Text;
6 import org.apache.hadoop.mapreduce.Job;
7 import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
8 import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
9
10 public class Driver4 {
11
12     public static void main(String[] args) throws Exception {
13         Configuration conf = new Configuration();
14         conf.set("m", "100");
15         conf.set("n", "100");
16         conf.set("p", "100");
17         Job job = Job.getInstance(conf, "JobName");
18         job.setJarByClass(Driver4.class);
19         // TODO: specify a mapper
20         job.setMapperClass(Mapper4.class);
21         // TODO: specify a reducer
22         job.setReducerClass(Reducer4.class);
23
24         // TODO: specify output types
25         job.setOutputKeyClass(Text.class);
26         job.setOutputValueClass(Text.class);
27
28         // TODO: specify input and output DIRECTORIES (not files)
29         FileInputFormat.setInputPaths(job, new Path("hdfs://localhost:9000/input/exp4_postlab"));
30         FileOutputFormat.setOutputPath(job, new Path("hdfs://localhost:9000/output/exp4_postlab"));
31
32         long start = System.nanoTime();
33         System.out.println("Execution started at: "+Long.toString(start)+" : "+java.time.LocalDateTime.now());
34
35         System.out.println("MapReduce job successful "+job.waitForCompletion(true));
36
37         long end = System.nanoTime();
38         System.out.println("Execution started at: "+Long.toString(end)+" : "+java.time.LocalDateTime.now());
39         long elapsed = end-start;
40         System.out.println("Time elapsed: "+Long.toString(elapsed)+"ns");
41     }
42
43 }
```

Execution started at: 1542792734619500 : 11:42:58.288

MapReduce job successful true

Execution started at: 1542801765783800 : 11:43:07.277

Time elapsed: 9031164300ns